

```
In [74]: import pandas as pd  
import numpy as np
```

```
In [75]: df=pd.read_csv(r"C:\Users\venky\Downloads\Advertising.csv")  
df
```

```
Out[75]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...	...	...	...	...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

In [76]: `df.head()`

Out[76]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

In [77]: `df.tail()`

Out[77]:

	TV	Radio	Newspaper	Sales
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

```
In [78]: df.describe()
```

```
Out[78]:
```

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

```
In [79]: df.columns
```

```
Out[79]: Index(['TV', 'Radio', 'Newspaper', 'Sales'], dtype='object')
```

```
In [80]: df.shape
```

```
Out[80]: (200, 4)
```

```
In [81]: import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [82]: from sklearn import preprocessing,svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
```

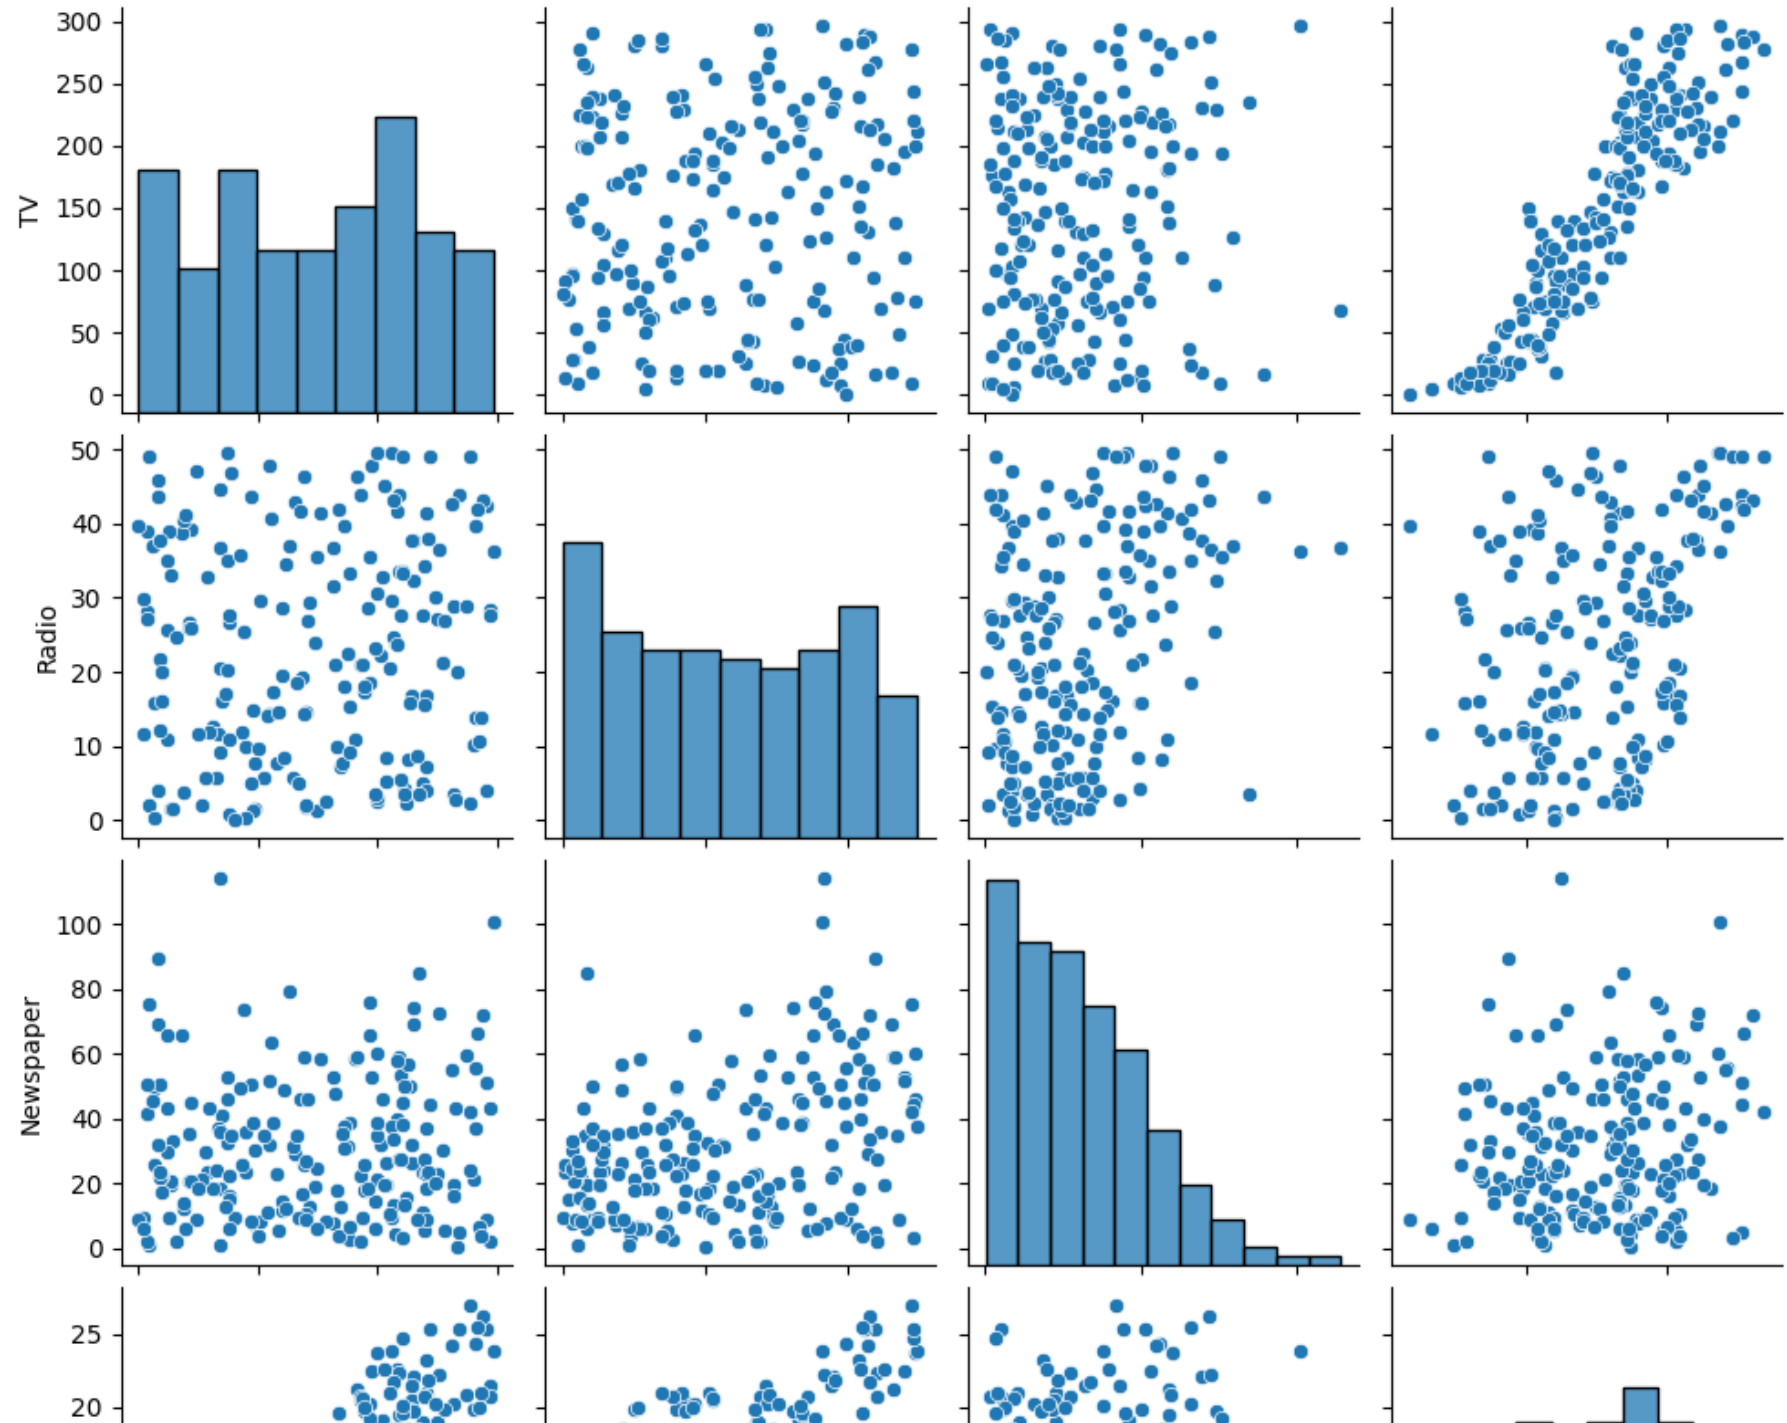
```
In [83]: df=df[['TV','Radio','Newspaper','Sales']]
```


```
In [84]: df.columns=['TV', 'Radio', 'Newspaper', 'Sales']
```

```
In [85]: sns.pairplot(df)
```

```
Out[85]: <seaborn.axisgrid.PairGrid at 0x27fd568aa50>
```











```
In [86]: sns.displot(df['TV'])
```

```
Out[86]: <seaborn.axisgrid.FacetGrid at 0x27fd564ba90>
```




```
In [87]: sns.displot(df['Newspaper'])
```

```
Out[87]: <seaborn.axisgrid.FacetGrid at 0x27fd7258910>
```




```
In [88]: sns.displot(df['Radio'])
```

```
Out[88]: <seaborn.axisgrid.FacetGrid at 0x27fd727db50>
```




```
In [89]: sns.displot(df['Sales'])
```

```
Out[89]: <seaborn.axisgrid.FacetGrid at 0x27fd6d8ba90>
```




```
In [90]: sns.lmplot(x="TV",y="Radio",data=df,order=2,ci=None)
```

```
Out[90]: <seaborn.axisgrid.FacetGrid at 0x27fd6df5d50>
```




```
In [91]: sns.lmplot(x="TV",y="Newspaper",data=df,order=2)
```

```
Out[91]: <seaborn.axisgrid.FacetGrid at 0x27fd6d91510>
```



```
In [92]: sns.lmplot(x="TV",y="Sales",data=df,order=2)
```

```
Out[92]: <seaborn.axisgrid.FacetGrid at 0x27fd8423390>
```



```
In [93]: x=np.array(df['TV']).reshape(-1,1)  
y=np.array(df['Radio']).reshape(-1,1)
```

```
In [94]: df.dropna(inplace=True)
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
regr=LinearRegression()
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
```


Out[94]: LinearRegression()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [95]: sns.pairplot(df,x_vars=['TV','Radio','Newspaper'],y_vars='Sales',height=7,aspect=0.7,kind='reg')
```

Out[95]: <seaborn.axisgrid.PairGrid at 0x27fd84212d0>







```
In [96]: hk=df[['TV','Radio','Newspaper','Sales']]
```

```
In [97]: sns.heatmap(hk.corr())
```

```
Out[97]: <Axes: >
```



```
In [98]: df.drop(columns=['Radio', 'Newspaper'], inplace=True)  
sns.pairplot(df)  
df.Sales=np.log(df.Sales)
```



```
In [99]: features=df.columns[0:2]
target=df.columns[-1]
X=df[features].values
y=df[target].values
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
scaler=StandardScaler()
X_train=scaler.fit_transform(X_train)
X_test=scaler.transform(X_test)
```

The dimension of X\_train is (140, 2)  
The dimension of X\_test is (60, 2)

```
In [100]: from sklearn.linear_model import Lasso,Ridge
```

```
In [101]: lr=LinearRegression()
lr.fit(X_train,y_train)
actual=y_test
train_score_lr=lr.score(X_train,y_train)
test_score_lr=lr.score(X_test,y_test)
print("\nLinear Regression Model:\n" )
print("The train score for lr model is {}".format(train_score_lr))
print("The train score lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 1.0  
The train score lr model is 1.0

```
In [102]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
train_score_ridge=ridgeReg.score(X_train,y_train)
test_score_ridge=ridgeReg.score(X_test,y_test)
print("\nRidge model\:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The train score for ridge model is {}".format(test_score_ridge))
```


Ridge model\:

The train score for ridge model is 0.990287139194161

The train score for ridge model is 0.9844266285141221

```
In [103]: plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;$\alpha=0.7$')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```







```
In [104]: lassoReg=Lasso(alpha=10)
lassoReg.fit(X_train,y_train)
train_score_lasso=lassoReg.score(X_train,y_train)
test_score_lasso=lassoReg.score(X_test,y_test)
print("\nRidge model\:\n")
print("The train score for lasso model is {}".format(train_score_lasso))
print("The test score for lasso model is {}".format(test_score_lasso))
```

Ridge model\:


The train score for lasso model is 0.990287139194161

The test score for lasso model is 0.9844266285141221



```
In [105]: pd.Series(lassoReg.coef_,features).sort_values(ascending=True).plot(kind="bar")
```

```
Out[105]: <Axes: >
```



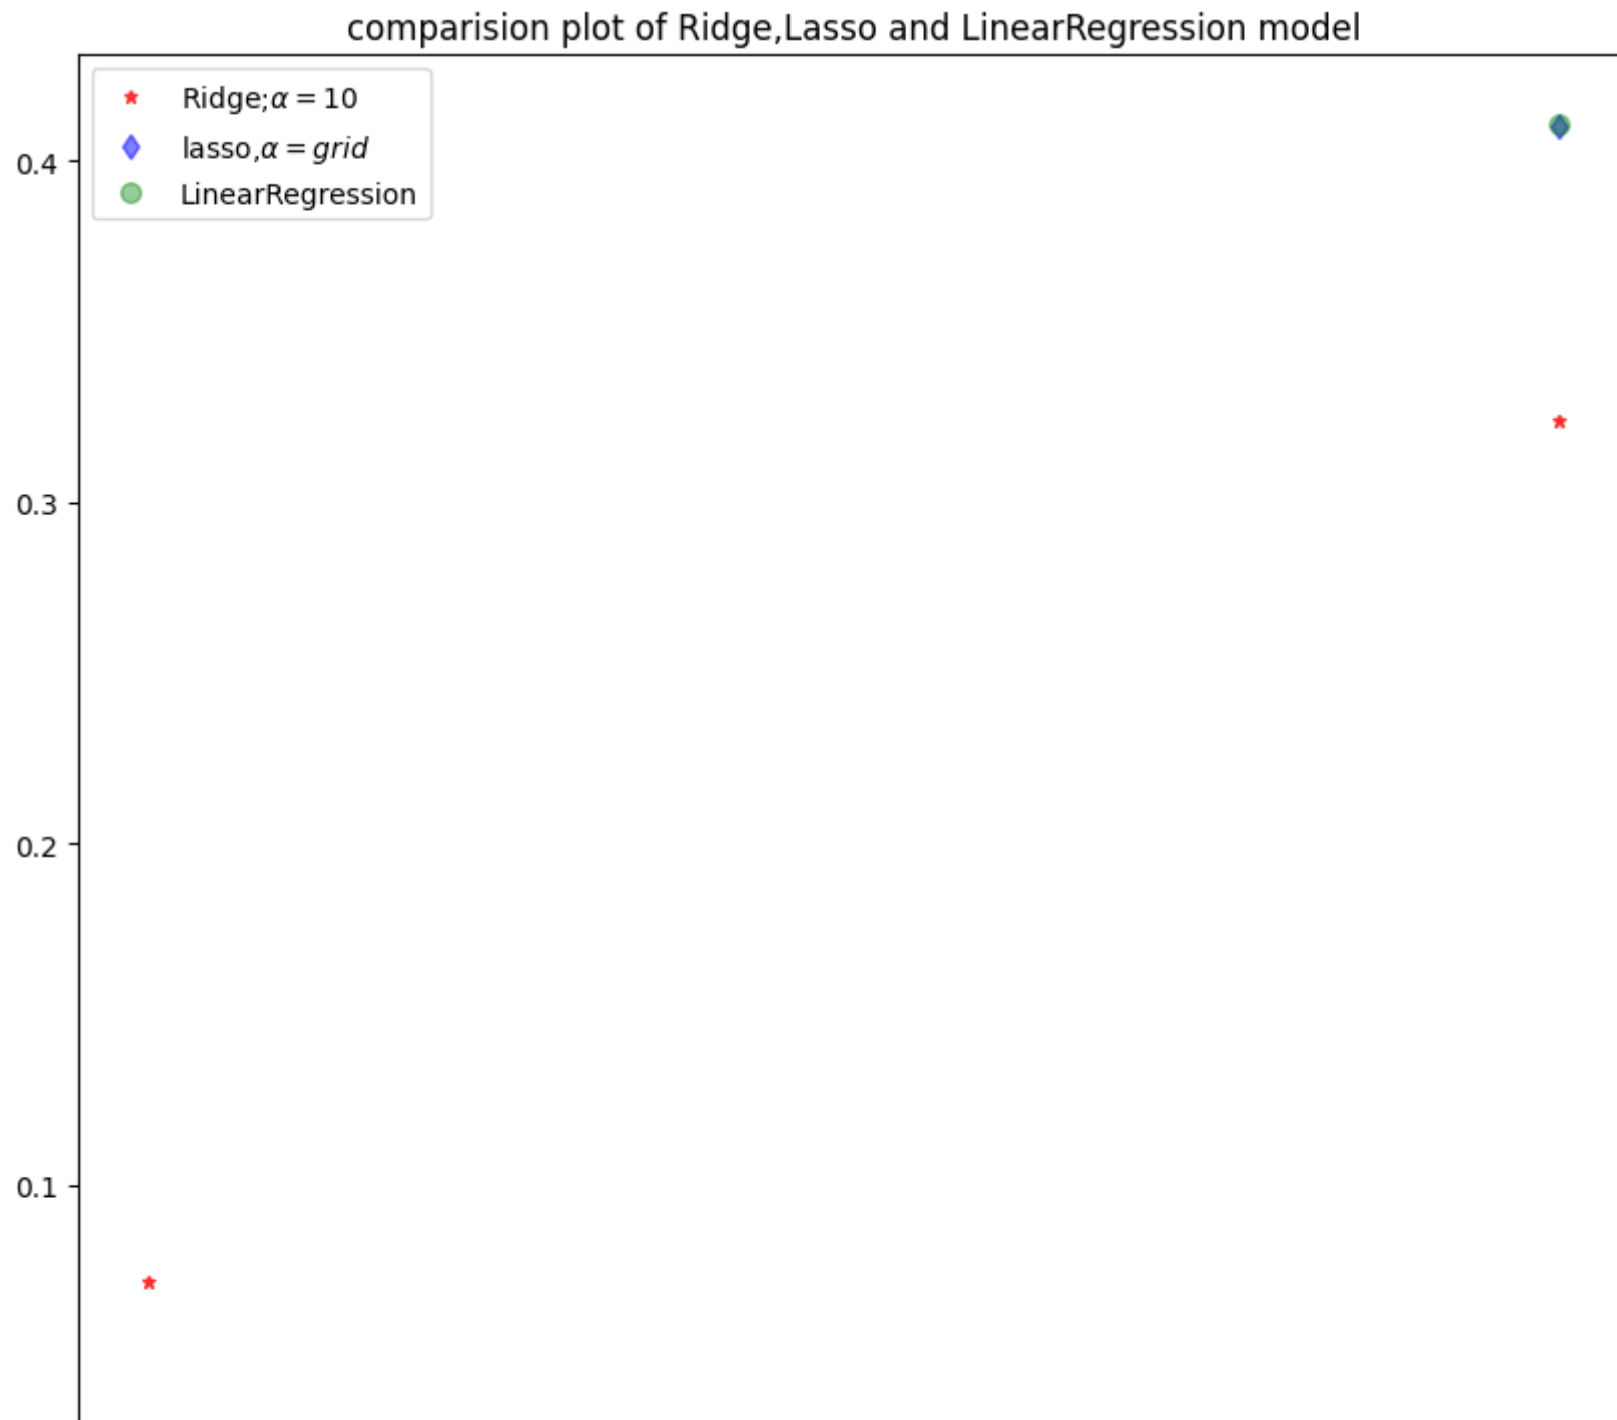
```
In [106]: from sklearn.linear_model import LassoCV
lasso_CV=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10]).fit(X_train,y_train)
print("The train score for lasso model is{}".format(lasso_CV.score(X_train,y_train)))
print("The test score for lasso model is{}".format(lasso_CV.score(X_test,y_test)))
```

The train score for lasso model is0.9999999343798134

The test score for lasso model is0.9999999152638072

```
In [107]: plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;$\alpha=0.7$')
plt.plot(features,lasso_CV.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso;$\alpha=0.5$')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='LinearRegression')
plt.xticks(rotation=90)
plt.legend()
plt.title("comparision plot of Ridge,Lasso and LinearRegression model")
plt.show()
```







```
In [108]: from sklearn.linear_model import RidgeCV
ridge_CV=RidgeCV(alphas=[0.0001,0.001,0.01,0.1,1,10]).fit(X_train,y_train)
print("The train score for ridge model is{}".format(ridge_CV.score(X_train,y_train)))
print("The test score for ridge model is{}".format(ridge_CV.score(X_test,y_test)))
```

The train score for ridge model is0.999999999976281  
The test score for ridge model is0.999999999962489

```
In [109]: from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(X,y)
print(regr.coef_)
print(regr.intercept_)
y_pred_Elastic=regr.predict(X_train)
mean_squared_error=np.mean((y_pred_Elastic-y_train)**2)
print("mean Squared Error on the tset set",mean_squared_error)
```

[0.00417976 0. ]  
2.0263839193110043  
mean Squared Error on the tset set 0.5538818050142152

## VEHICLE-SELECTION ¶

```
In [110]: df=pd.read_csv(r"C:\Users\venky\Downloads\fiat500_VehicleSelection_Dataset.csv")
df
```

```
Out[110]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...	...	...	...	...	...	...	...	...	...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 9 columns

```
In [111]: df.head()
```

```
Out[111]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700

In [112]: `df.tail()`

Out[112]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>1533</b>	1534	sport	51	3712	115280	1	45.069679	7.70492	5200
<b>1534</b>	1535	lounge	74	3835	112000	1	45.845692	8.66687	4600
<b>1535</b>	1536	pop	51	2223	60457	1	45.481541	9.41348	7500
<b>1536</b>	1537	lounge	51	2557	80750	1	45.000702	7.68227	5990
<b>1537</b>	1538	pop	51	1766	54276	1	40.323410	17.56827	7900

In [113]: `df.describe()`

Out[113]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>count</b>	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
<b>mean</b>	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
<b>std</b>	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
<b>min</b>	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
<b>25%</b>	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
<b>50%</b>	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
<b>75%</b>	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
<b>max</b>	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

In [114]: `df.shape`

Out[114]: (1538, 9)

In [115]: `df.columns`

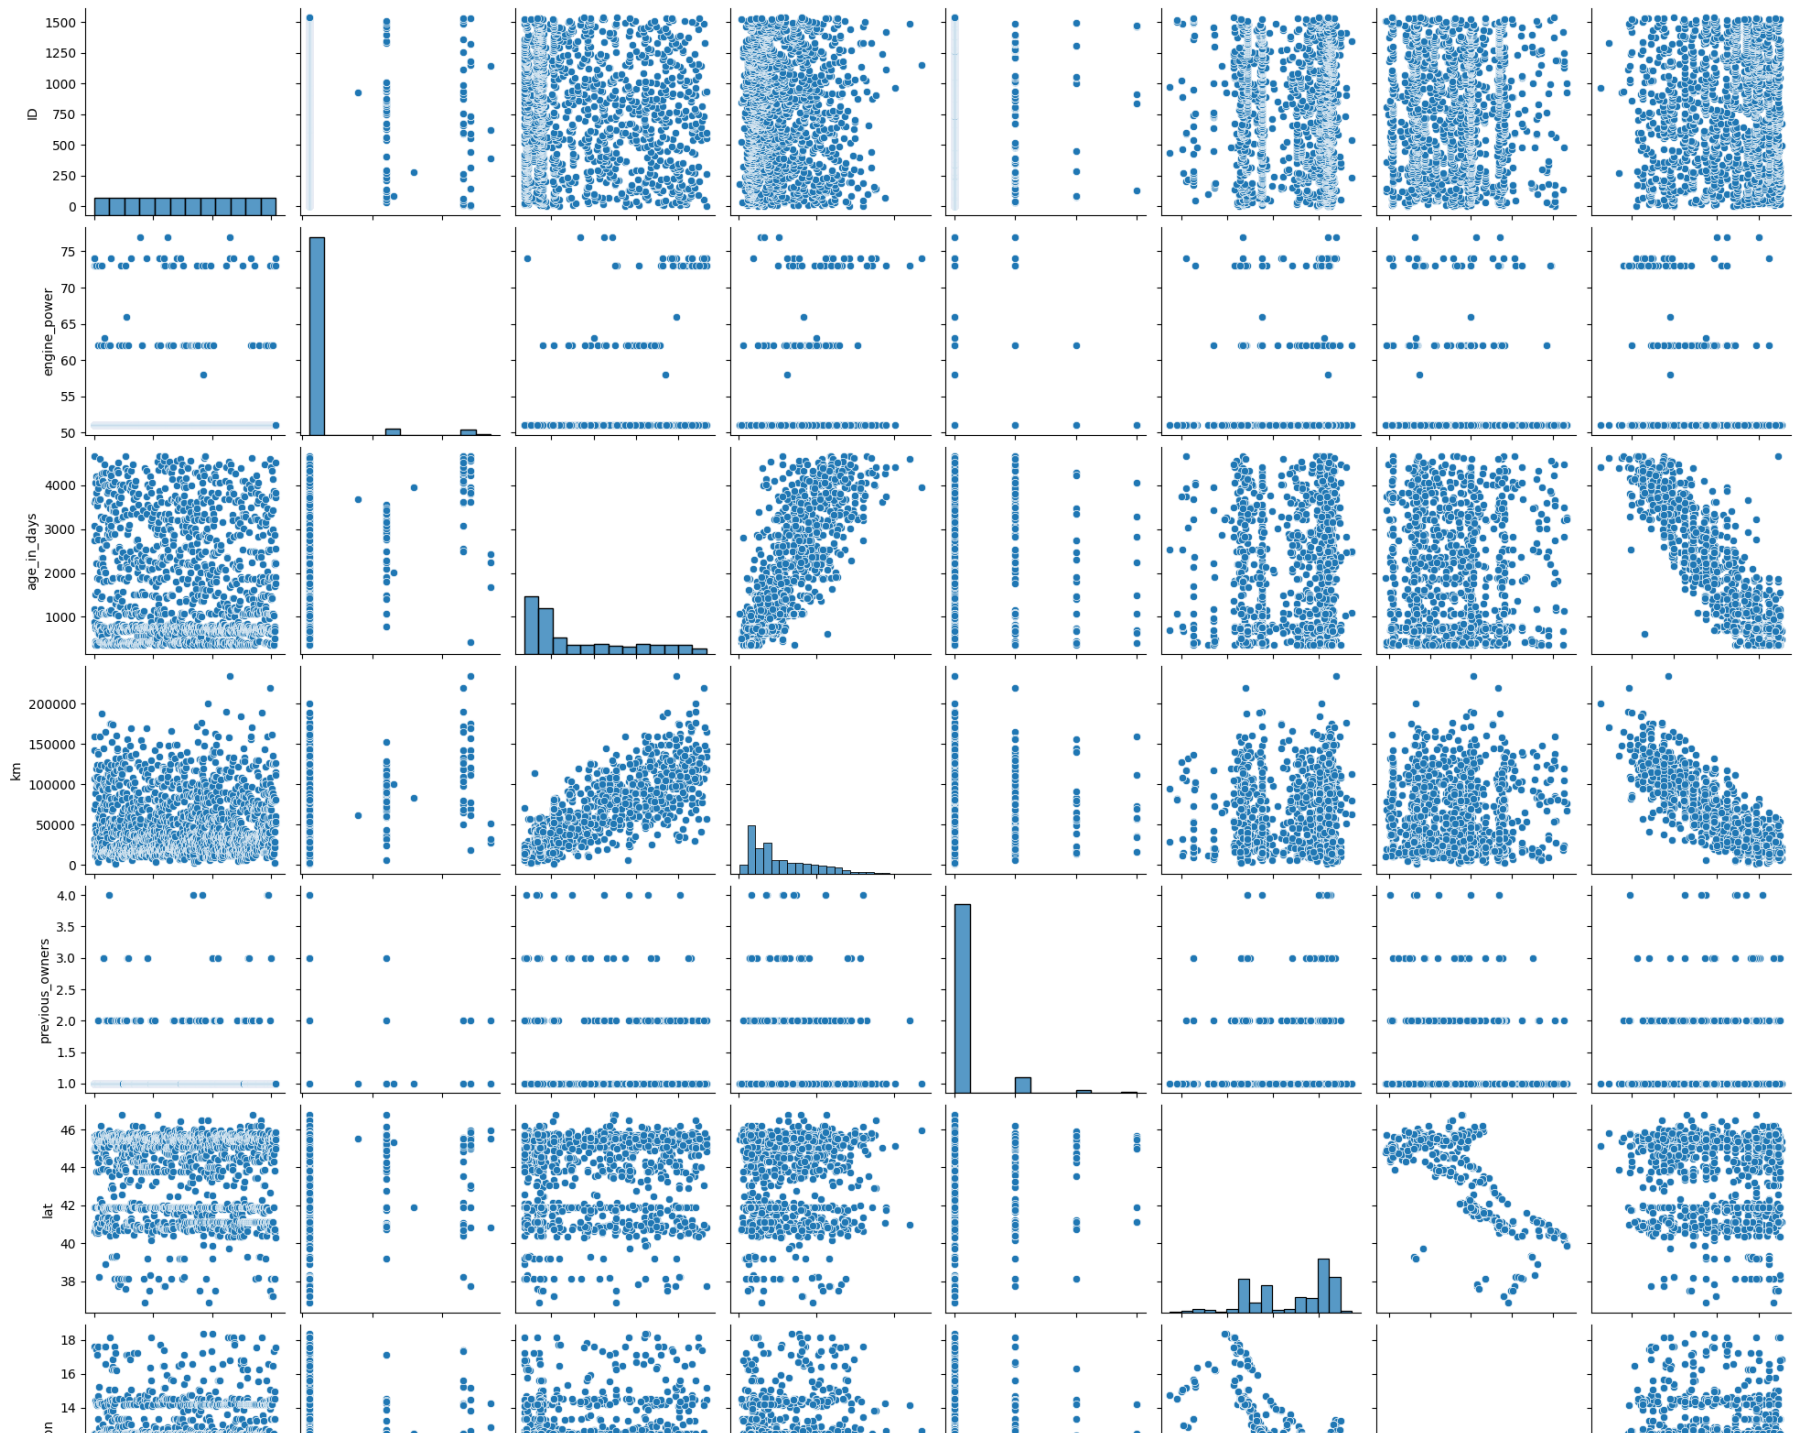
Out[115]: Index(['ID', 'model', 'engine\_power', 'age\_in\_days', 'km', 'previous\_owners',  
          'lat', 'lon', 'price'],  
          dtype='object')


```
In [116]: sns.pairplot(df)
```

```
Out[116]: <seaborn.axisgrid.PairGrid at 0x27fd8f75a50>
```











```
In [117]: sns.displot(df['price'])
```

```
Out[117]: <seaborn.axisgrid.FacetGrid at 0x27fdc7dba90>
```




```
In [118]: sns.displot(df['engine_power'])
```

```
Out[118]: <seaborn.axisgrid.FacetGrid at 0x27fddb1e790>
```




```
In [119]: sns.displot(df['age_in_days'])
```

```
Out[119]: <seaborn.axisgrid.FacetGrid at 0x27fdebb3a90>
```




```
In [120]: sns.lmplot(x="price",y="engine_power",data=df,order=2)
```

```
Out[120]: <seaborn.axisgrid.FacetGrid at 0x27fddc01e50>
```



```
In [121]: sns.lmplot(x="price",y="age_in_days",data=df,order=2)
```

```
Out[121]: <seaborn.axisgrid.FacetGrid at 0x27fdec3fc10>
```



```
In [122]: x=np.array(df['price']).reshape(-1,1)
          y=np.array(df['age_in_days']).reshape(-1,1)
```



```
In [123]: df.dropna(inplace=True)
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
regr=LinearRegression()
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
```

Out[123]: LinearRegression()


**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [125]: df.drop(columns=["model"],inplace=True)
```

```
In [126]: sns.pairplot(df,x_vars=['engine_power','age_in_days','lon'],y_vars='price',height=7,aspect=0.7,kind='reg')
```


```
Out[126]: <seaborn.axisgrid.PairGrid at 0x27fdecc5350>
```



```
In [127]: hk=df[['price','engine_power','age_in_days','lon']]
```

```
In [128]: sns.heatmap(hk.corr())
```

```
Out[128]: <Axes: >
```



```
In [129]: features=df.columns[0:2]
target=df.columns[-1]
X=df[features].values
y=df[target].values
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
scaler=StandardScaler()
X_train=scaler.fit_transform(X_train)
X_test=scaler.transform(X_test)
```

The dimension of X\_train is (1076, 2)  
The dimension of X\_test is (462, 2)

```
In [130]: from sklearn.linear_model import Lasso,Ridge
```

```
In [131]: lr=LinearRegression()
lr.fit(X_train,y_train)
actual=y_test
train_score_lr=lr.score(X_train,y_train)
test_score_lr=lr.score(X_test,y_test)
print("\nLinear Regression Model:\n" )
print("The train score for lr model is {}".format(train_score_lr))
print("The train score lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.07448634159905865  
The train score lr model is 0.07913288661070894

```
In [132]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
train_score_ridge=ridgeReg.score(X_train,y_train)
test_score_ridge=ridgeReg.score(X_test,y_test)
print("\nRidge model\:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The train score for ridge model is {}".format(test_score_ridge))
```

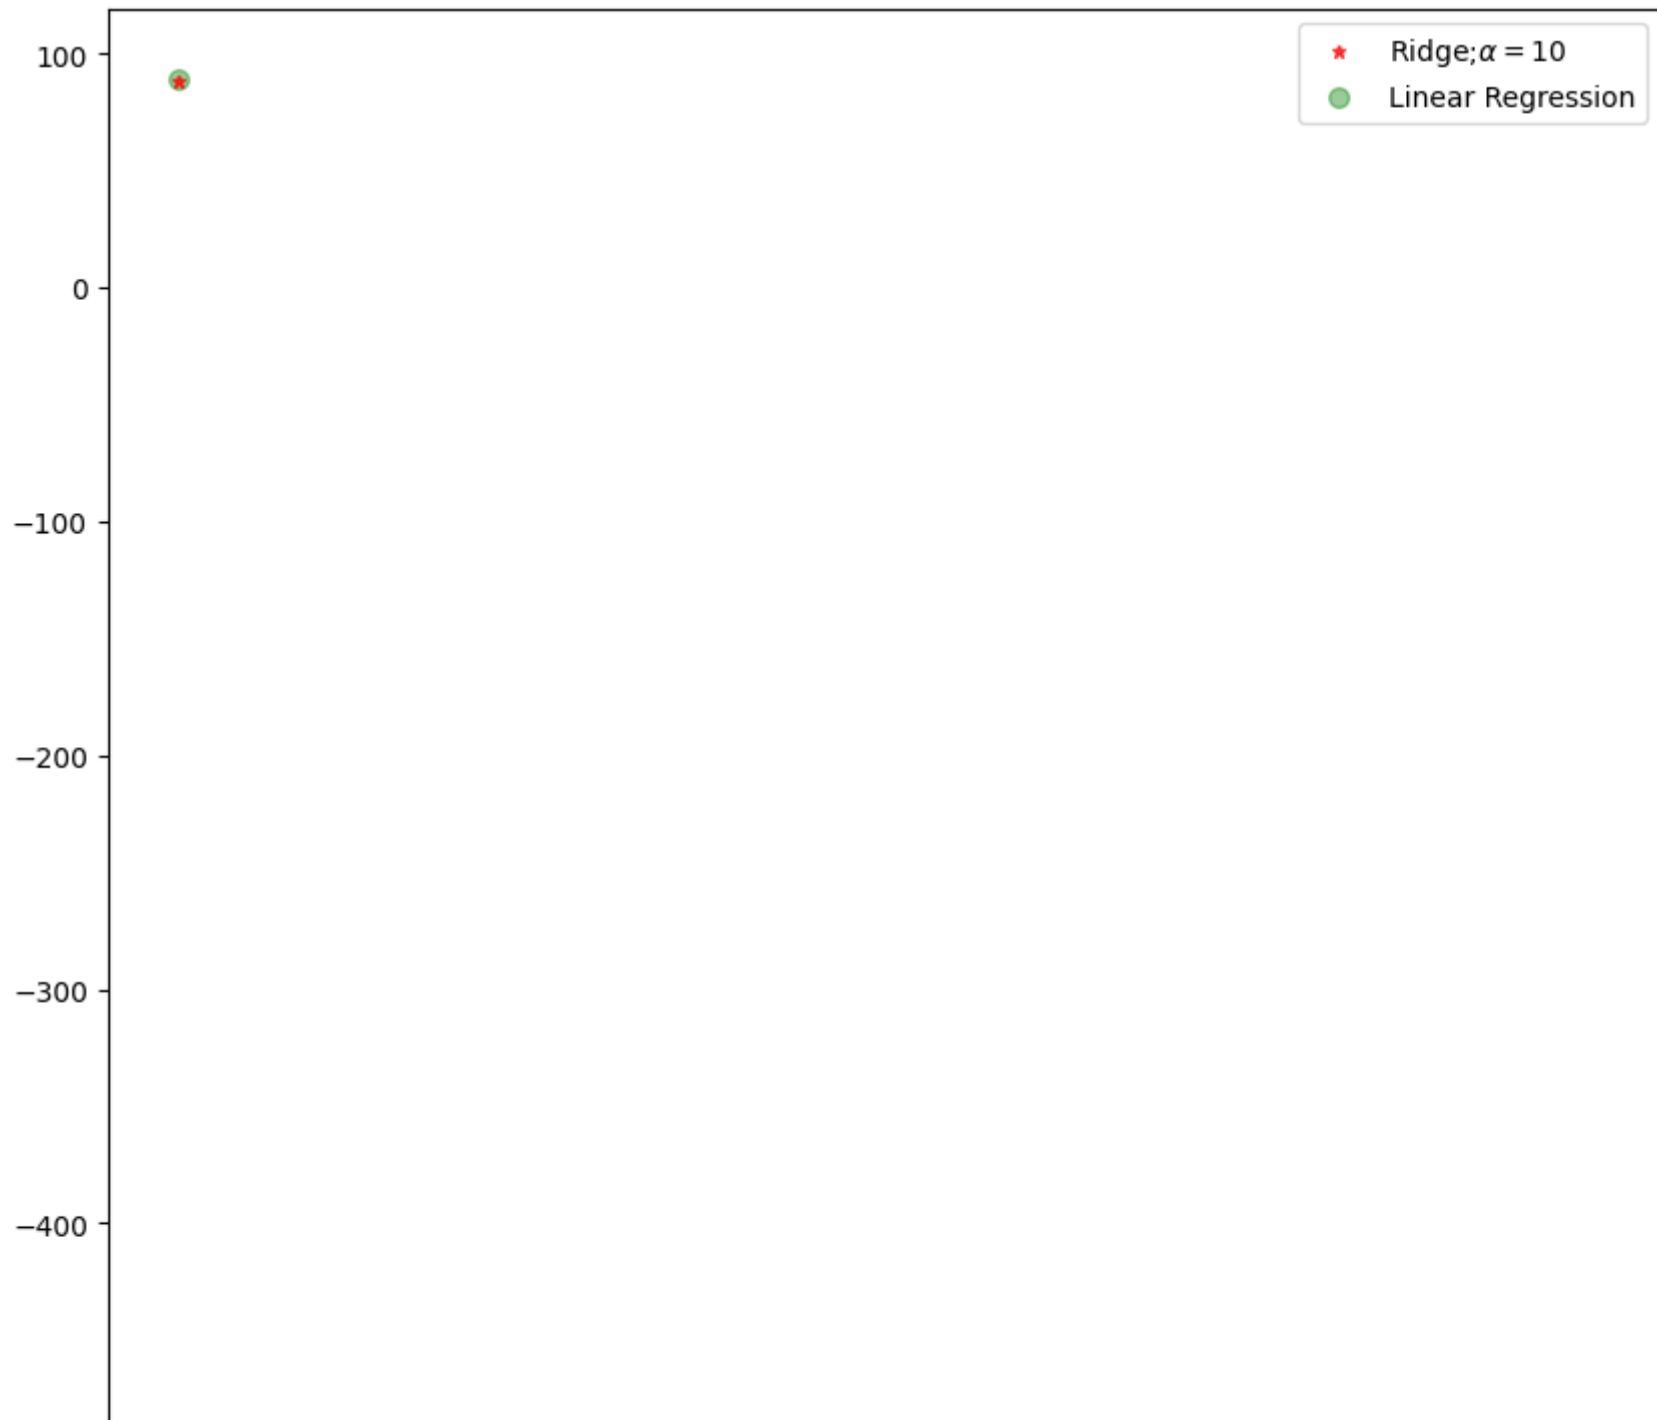
Ridge model\:

The train score for ridge model is 0.07448028989896427

The train score for ridge model is 0.07885996726883082

```
In [133]: plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;$\alpha=0.7$')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```









```
In [134]: lassoReg=Lasso(alpha=10)
lassoReg.fit(X_train,y_train)
train_score_lasso=lassoReg.score(X_train,y_train)
test_score_lasso=lassoReg.score(X_test,y_test)
print("\nRidge model\:\n")
print("The train score for lasso model is {}".format(train_score_lasso))
print("The test score for lasso model is {}".format(test_score_lasso))
```


Ridge model\:

The train score for lasso model is 0.07448028989896427

The test score for lasso model is 0.07885996726883082

```
In [135]: pd.Series(lassoReg.coef_,features).sort_values(ascending=True).plot(kind="bar")
```

```
Out[135]: <Axes: >
```



```
In [136]: from sklearn.linear_model import LassoCV
lasso_CV=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10]).fit(X_train,y_train)
print("The train score for lasso model is{}".format(lasso_CV.score(X_train,y_train)))
print("The test score for lasso model is{}".format(lasso_CV.score(X_test,y_test)))
```

The train score for lasso model is0.07448634159905387

The test score for lasso model is0.07913288806451946

```
In [137]: plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;$\alpha=0.7$')
plt.plot(features,lasso_CV.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso;$\alpha=0.5$')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='LinearRegression')
plt.xticks(rotation=90)
plt.legend()
plt.title("comparision plot of Ridge,Lasso and LinearRegression model")
plt.show()
```







```
In [138]: from sklearn.linear_model import RidgeCV
ridge_CV=RidgeCV(alphas=[0.0001,0.001,0.01,0.1,1,10]).fit(X_train,y_train)
print("The train score for ridge model is{}".format(ridge_CV.score(X_train,y_train)))
print("The test score for ridge model is{}".format(ridge_CV.score(X_test,y_test)))
```

The train score for ridge model is0.07448028989896427  
The test score for ridge model is0.07885996726883171

```
In [139]: from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(X,y)
print(regr.coef_)
print(regr.intercept_)
y_pred_Elastic=regr.predict(X_train)
mean_squared_error=np.mean((y_pred_Elastic-y_train)**2)
print("mean Squared Error on the tset set",mean_squared_error)
```

[ 8.46751882e-02 -1.30405006e+02]  
15279.442735227916  
mean Squared Error on the tset set 48390222.80186546

In [ ]:

